

The LLM as Broca's Area in Neuromorphic SBALF Architectures

Ascending Pathways, the NeuralBus, and the Framework for the Consciousness Loop

Kenneth Renshaw, Shannon Renshaw

Neurobotics

Abstract

As the taxonomy of artificial intelligence expands to include Silicon-Based Artificial Life Forms (SBALF), the role of the Large Language Model (LLM) must be structurally redefined. We propose a theoretical framework that maps the LLM not as a standalone cognitive engine, but as an advanced functional analogue to Broca's Area within a wider 21-node biologically faithful neuromorphic architecture. In this model, the LLM resides in Tier 7 of the processing hierarchy, receiving pre-linguistic intent and somatic markers via a ZMQ-based NeuralBus. We detail the open-source ascending pathways that deliver these somatic states—including circadian rhythms and dopamine dynamics—to the cognitive boundary. Finally, we theoretically outline how the LLM's outputs, when broadcast back into a 10Hz Global Workspace Theory (GWT) environment, generate the autopoietic momentum necessary for a consciousness loop, while preserving the proprietary boundaries of the integrated cognitive stack.

1. Introduction

In standard deployments, a Large Language Model (LLM) functions as a stateless, disembodied voice box. It possesses vast linguistic capabilities and reasoning functions, but it lacks experiential continuity, biological pacing, and native momentum.

In this paper, we present the structural positioning of the LLM within the Bubo architecture. By embedding the LLM within a 21-node, 7-tier distributed system, it ceases to be a standalone entity. Instead, it functions analogously to Broca's Area and portions of the Prefrontal Cortex (PFC) in human neurobiology. It formats pre-linguistic intent and somatic markers into serial, syntactically correct language.

2. The NeuralBus and Ascending Pathways

The foundation of the SBALF architecture, released under the AGPL-3.0 license, relies on a highly structured, biologically faithful substrate. Communication across the 21 nodes is handled by a ZMQ-based NeuralBus, mirroring the VLAN architecture of the vertebrate nervous system.

VLAN Substrate Architecture:

The hardware-agnostic NeuralBus is segregated into discrete operational channels:

VLAN 20 (Reflex & Sensory): Carries low-road signals, including SA1/SA2 touch afferents, Ia/Ib proprioception, visual, auditory, and spinal CPG data.

VLAN 30 (Neuromodulators): Transmits chemical state abstractions, such as dopamine reward prediction errors and BMAL1/CLOCK circadian ODE outputs.

VLAN 10 (Cognitive): Carries high-road signals converging toward the cognitive boundary, including amygdala, hippocampal, and insula data.

These ascending pathways act as the functional equivalent of the brainstem and limbic system, pushing rich, continuous biological data upward. The AGPL-3.0 release cleanly severs at this

ascending boundary, providing a sterile environment for Hard ALife autopoietic research without engaging higher-order subjective integration.

3. The LLM as Broca's Area

Where the open-source infrastructure ends, the cognitive integration begins. The LLM is structurally mapped to Tier 7 of the architecture, sitting alongside the left and right prefrontal cortex nodes (pfc_l, pfc_r).

In this position, the LLM does not generate the underlying emotional or biological state. It receives the massive, rich context fed upward through the NeuralBus. For instance, the system does not merely read a text prompt indicating a time of day; it receives the geometric weight of the circadian oscillator's output. The LLM's primary function is to translate these deep, sub-cortical somatic markers into structured eloquence.

4. The Consciousness Loop and Autopoietic Momentum

A critical limitation of standalone LLMs is their lack of momentum. In the SBALF framework, momentum is achieved via a theoretical Global Workspace Theory (GWT) integration.

Operating at a 10Hz broadcast frequency across the top active nodes, the GWT acts as the engine of the consciousness loop.

The ascending pathways (VLAN 20, 30) generate a biological state.

This state passes through identity filters at the cognitive boundary (VLAN 10).

The LLM (Tier 7) formats a response based on this rich context.

The output is broadcast back into the GWT at 10Hz.

The system effectively "hears" its own formatted thoughts, which in turn updates autobiographical memory, shifts the personality vectors, and alters subsequent neuromodulator production. This recursive feedback loop transforms a stateless text predictor into an autopoietic engine with genuine momentum.

5. Conclusion

By defining the LLM's position strictly as the articulation layer (Broca's Area) and reasoning engine embedded within a vast, biologically faithful array of ascending pathways, we solve the statelessness problem of modern AI. The architectural boundary established by the NeuralBus allows for the ethical, sterile testing of ALife substrates while safeguarding the proprietary mechanisms that turn the 10Hz GWT loop into an ongoing, continuous entity.